

FunSearch & AlphaEvolve

Bartosz Piotrowski

August 21, 2025

The Reasoning Reading Group @ FAIR

Motivation

Preconditions:

- A problem of finding optimal heuristic / program / function.
- A pre-trained, coding-capable LLM.
- An automated evaluator returning a scalar score.

Motivation

Preconditions:

- A problem of finding optimal heuristic / program / function.
- A pre-trained, coding-capable LLM.
- An automated evaluator returning a scalar score.

The simplest strategy:

- generating multiple independent samples from the LLM with non-zero temperature,
- evaluating all of them,
- selecting the best one.

Motivation

Preconditions:

- A problem of finding optimal heuristic / program / function.
- A pre-trained, coding-capable LLM.
- An automated evaluator returning a scalar score.

The simplest strategy:

- generating multiple independent samples from the LLM with non-zero temperature,
- evaluating all of them,
- selecting the best one.

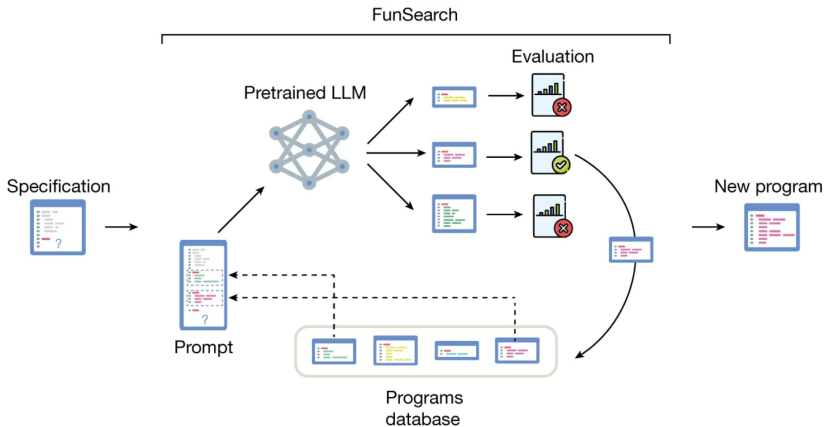
But instead of sampling independently, can we **incorporate the evaluator feedback** into subsequent generations?

FunSearch

Key ingredients:

- *best-shot* prompting,
- a growing database of programs,
- an evolutionary strategy acting on it.

FunSearch



Database of programs

- Several islands/subpopulations growing independently.
- Higher-scoring programs, but also shorter ones, are prioritized.
- Less good programs are eventually discarded.
- Different islands are mixed with each other to an extent.
- Why multiple islands? For diversity.

Best-shot prompting

- k good programs per prompt sampled, for each island.
- Information which one is better incorporated into the prompt (v_0, v_1, \dots) .
- In the actual experiments, $k = 2$.

Other details

- **Model:**
 - Codey, a PaLM2 model fine-tuned on code.
 - Smaller, faster-inference model chosen.
- Implementation with **three asynchronous workers:**
 - database,
 - generator,
 - evaluator.
- Prompting with **templates / skeletons of programs.**
 - The LLM asked to modify only the essential function.

Problems tackled

Two known combinatorics problems:


- Cap set problem.
- (Online) bin packing.

Cap set problem

What is the largest possible set of vectors in \mathbb{Z}_3^n (*cap set*) such that no three vectors sum to zero?


Cap set problem

What is the largest possible set of vectors in \mathbb{Z}_3^n (*cap set*) such that no three vectors sum to zero?




Cap set problem

What is the largest possible set of vectors in \mathbb{Z}_3^n (*cap set*) such that no three vectors sum to zero?



Cap set problem

What is the largest possible set of vectors in \mathbb{Z}_3^n (*cap set*) such that no three vectors sum to zero?



n	3	4	5	6	7	8
Best known	9	20	45	112	236	496
FunSearch	9	20	45	112	236	512

Cap set problem template

```
"""Finds large cap sets."""
import numpy as np
import utils_capset

# Function to be executed by FunSearch.
def main(n):
    """Runs `solve` on `n`-dimensional cap set and
    ↳ evaluates the output."""
    solution = solve(n)
    return evaluate(solution, n)

def evaluate(candidate_set, n):
    """Returns size of candidate_set if it is a cap
    ↳ set, None otherwise."""
    if utils_capset.is_capset(candidate_set, n):
        return len(candidate_set)
    else:
        return None

def solve(n):
    """Builds a cap set of dimension `n` using
    ↳ `priority` function."""
    # Precompute all priority scores.
    elements = utils_capset.get_all_elements(n)
    scores = [priority(el, n) for el in elements]
    # Sort elements according to the scores.
    elements = elements[np.argsort(scores,
    ↳ kind='stable')[:-1]]

    # Build `capset` greedily, using scores for
    ↳ prioritization.
    capset = []
    for element in elements:
        if utils_capset.can_be_added(element, capset):
            capset.append(element)
    return capset

# Function to be evolved by FunSearch.
def priority(element, n):
    """Returns the priority with which we want to add
    ↳ `element` to the cap set."""
    return 0.0
```

Cap set solution – the priority function

```
def priority(el: tuple[int, ...],
    ↪ n: int) -> float:
    score = n
    in_el = 0
    el_count = el.count(0)

    if el_count == 0:
        score += n**2
        if el[1] == el[-1]:
            score *= 1.5
        if el[2] == el[-2]:
            score *= 1.5
        if el[3] == el[-3]:
            score *= 1.5
    else:
        if el[1] == el[-1]:
            score *= 0.5
        if el[2] == el[-2]:
            score *= 0.5

    for e in el:
        if e == 0:
            if in_el == 0:
                score *= n * 0.5
            elif in_el == el_count - 1:
                score *= 0.5
            else:
                score *= n * 0.5 ** in_el
            in_el += 1
        else:
            score += 1

    if el[1] == el[-1]:
        score *= 1.5
    if el[2] == el[-2]:
        score *= 1.5

    return score
```


AlphaEvolve

tldr: FunSearch scaled-up in multiple dimensions.

<i>FunSearch</i> [83]	<i>AlphaEvolve</i>
evolves single function	evolves entire code file
evolves up to 10-20 lines of code	evolves up to hundreds of lines of code
evolves code in Python	evolves any language
needs fast evaluation (≤ 20 min on 1 CPU)	can evaluate for hours, in parallel, on accelerators
millions of LLM samples used	thousands of LLM samples suffice
small LLMs used; no benefit from larger	benefits from SOTA LLMs
minimal context (only previous solutions)	rich context and feedback in prompts
optimizes single metric	can simultaneously optimize multiple metrics

AlphaEvolve – notable differences

- Larger base LLMs (Gemini 2.0 Flash, Gemini 2.0 Pro).
- Much richer prompts:
 - context for the problem provided,
 - detailed evaluation results included,
 - meta prompt evolution.
- Output format: a **diff** rather than a new program version.
- Evaluation:
 - many evaluators,
 - including LLM judges.
- Improved evolution strategy:
 - different islands (as before),
 - MAP elites algorithm.

Problems tackled

50 mathematical problems:

- Fast matrix multiplication, kissing number problem, ...
- For 75% problems AlphaEvolve matched the optimal known solution.
- For 20% problems AlphaEvolve found better solution.

Also practical problems related to computational infrastructure:


- Scheduling jobs on a cluster, optimizing JAX kernel, ...
- Some of the solutions found by AlphaEvolve went into production at Google.

Optimizing matrix multiplications

- The simple algorithm for multiplying $n \times m$ and $m \times k$ matrices requires nmk scalar multiplications.
- But Volker Strassen in 1969 showed it can be done with fewer multiplications.
- The optimal algorithm in general not known.
- AlphaEvolve found improvements for multiple specific cases.

$\langle m, n, p \rangle$	best known [reference]	<i>AlphaEvolve</i>
$\langle 2, 4, 5 \rangle$	33 [42]	32
$\langle 2, 4, 7 \rangle$	46 [93]	45
$\langle 2, 4, 8 \rangle$	52 [93]	51
$\langle 2, 5, 6 \rangle$	48 [93]	47
$\langle 3, 3, 3 \rangle$	23 [52]	23
$\langle 3, 4, 6 \rangle$	56 [48]	54
$\langle 3, 4, 7 \rangle$	66 [91]	63
$\langle 3, 4, 8 \rangle$	75 [91]	74
$\langle 3, 5, 6 \rangle$	70 [48]	68
$\langle 3, 5, 7 \rangle$	82 [91]	80
$\langle 4, 4, 4 \rangle$	49 [95]	48
$\langle 4, 4, 5 \rangle$	62 [47]	61
$\langle 4, 4, 7 \rangle$	87 [93]	85
$\langle 4, 4, 8 \rangle$	98 [95]	96
$\langle 4, 5, 6 \rangle$	93 [48]	90
$\langle 5, 5, 5 \rangle$	93 [72]	93

Ablations



Closing remarks

- FunSearch/AlphaEvolve can be seen as:
 - mechanizing effective interaction with the LLM,
 - complex meta-generation framework.¹
- An interesting trade-off:
 - smaller, less clever, but faster LLM = more samples (FunSearch) vs
 - larger, more clever, but slower LLM = higher-quality samples (AlphaEvolve).
- Is it possible to adapt the AlphaEvolve-style framework to formal proving?
 - A notion of *partial proof progress* would be needed.
 - But isn't it similar to what Seed Proved does at test time?
- AlphaEvolve could be distilled into a better base LLM, which would result in *meta-evolution*.
- Open-source implementation: OpenEvolve.

¹See a nice NeurIPS tutorial: *Beyond Decoding: Meta-Generation Algorithms for Large Language Models*